

Activity_Course 6 Automatidata project lab

October 31, 2025

1 Automatidata project

Course 6 - The Nuts and bolts of machine learning

You are a data professional in a data analytics firm called Automatidata. Their client, the New York City Taxi & Limousine Commission (New York City TLC), was impressed with the work you have done and has requested that you build a machine learning model to predict if a customer will not leave a tip. They want to use the model in an app that will alert taxi drivers to customers who are unlikely to tip, since drivers depend on tips.

A notebook was structured and prepared to help you in this project. Please complete the following questions.

2 Course 6 End-of-course project: Build a machine learning model

In this activity, you will practice using tree-based modeling techniques to predict on a binary target class.

The purpose of this model is to find ways to generate more revenue for taxi cab drivers.

The goal of this model is to predict whether or not a customer is a generous tipper.

This activity has three parts:

Part 1: Ethical considerations * Consider the ethical implications of the request

- Should the objective of the model be adjusted?

Part 2: Feature engineering

- Perform feature selection, extraction, and transformation to prepare the data for modeling

Part 3: Modeling

- Build the models, evaluate them, and advise on next steps

Follow the instructions and answer the questions below to complete the activity. Then, complete an Executive Summary using the questions listed on the PACE Strategy Document.

Be sure to complete this activity before moving on. The next course item will provide you with a completed exemplar to compare to your own work.

3 Build a machine learning model

4 PACE stages

Throughout these project notebooks, you'll see references to the problem-solving framework PACE. The following notebook components are labeled with the respective PACE stage: Plan, Analyze, Construct, and Execute.

4.1 PACE: Plan

Consider the questions in your PACE Strategy Document to reflect on the Plan stage.

In this stage, consider the following questions:

1. What are you being asked to do?
2. What are the ethical implications of the model? What are the consequences of your model making errors?
 - What is the likely effect of the model when it predicts a false negative (i.e., when the model says a customer will give a tip, but they actually won't)?
 - What is the likely effect of the model when it predicts a false positive (i.e., when the model says a customer will not give a tip, but they actually will)?
3. Do the benefits of such a model outweigh the potential problems?
4. Would you proceed with the request to build this model? Why or why not?
5. Can the objective be modified to make it less problematic?

What are you being asked to do? You are being asked to build a machine learning model to predict whether a taxi customer will not leave a tip. The goal is to use this model to alert drivers about customers who are unlikely to tip so they can adjust their service or expectations.

What are the ethical implications of the model? What are the consequences of making errors?

False negatives (predicting a customer will tip when they actually won't) could lead to unfair driver expectations.

False positives (predicting a customer won't tip when they actually do) could lead to discriminatory behavior toward certain customers.

There's a risk of bias if the model indirectly correlates tipping behavior with sensitive factors such as location, time, or customer demographics.

Do the benefits of such a model outweigh the potential problems? Not necessarily. While it might increase revenue for drivers, it could also promote unfair or unethical treatment of passengers. The risk of reinforcing bias or discrimination is high if the model isn't carefully monitored and explained.

Would you proceed with the request to build this model? Why or why not? I would proceed only with modifications — specifically, I'd ensure that sensitive or proxy demographic variables are excluded, that outputs are anonymized, and that the model's purpose is reframed toward improving overall service quality, not labeling individuals.

Can the objective be modified to make it less problematic? Yes. Instead of predicting who won't tip, the objective could focus on predicting factors that influence tipping behavior or identifying service improvements that lead to higher tips. This makes the model more ethical and action-oriented.

Suppose you were to modify the modeling objective so, instead of predicting people who won't tip at all, you predicted people who are particularly generous—those who will tip 20% or more? Consider the following questions:

1. What features do you need to make this prediction?
2. What would be the target variable?
3. What metric should you use to evaluate your model? Do you have enough information to decide this now?

New objective: Predict which customers are particularly generous — e.g., those who tip 20% or more.

What features do you need to make this prediction?

Trip distance, duration, fare amount, payment method, pickup/drop-off time, location, and ride type.

Possibly contextual data like weather or day of the week.

What would be the target variable? A binary variable:

1 = customer tips 20% or more

0 = customer tips less than 20%

What metric should you use to evaluate your model?

Use F1 score to balance precision and recall, since both false positives and false negatives matter.

Also review ROC-AUC for overall discriminative ability.

At this stage, there may not yet be enough information to finalize the metric, but F1 is a good starting point for an imbalanced dataset.

Complete the following steps to begin:

4.1.1 Task 1. Imports and data loading

Import packages and libraries needed to build and evaluate random forest and XGBoost classification models.

```
[197]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.metrics import accuracy_score, precision_score, recall_score,
f1_score, confusion_matrix, ConfusionMatrixDisplay, RocCurveDisplay
```

```
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

# This is the function that helps plot feature importance
from xgboost import plot_importance
```

```
[123]: # RUN THIS CELL TO SEE ALL COLUMNS
# This lets us see all of the columns, preventing Jupyter from redacting them.
pd.set_option('display.max_columns', None)
```

Begin by reading in the data. There are two dataframes: one containing the original data, the other containing the mean durations, mean distances, and predicted fares from the previous course's project called nyc_preds_means.csv.

Note: Pandas reads in the dataset as df0, now inspect the first five rows. As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed with this lab. Please continue with this activity by completing the following instructions.

```
[124]: # RUN THE CELL BELOW TO IMPORT YOUR DATA.
```

```
# Load dataset into dataframe
df0 = pd.read_csv('2017_Yellow_Trip_Data.csv')

# Import predicted fares and mean distance and duration from previous course
nyc_preds_means = pd.read_csv('nyc_preds_means.csv')
```

Inspect the first few rows of df0.

```
[125]: # Inspect the first few rows of df0
df0.head(10)
```

```
[125]:   Unnamed: 0  VendorID      tpep_pickup_datetime      tpep_dropoff_datetime
  passenger_count  trip_distance  RatecodeID \
0            24870114          2  03/25/2017 8:55:43 AM  03/25/2017 9:09:47 AM
6             3.34          1
1            35634249          1  04/11/2017 2:53:28 PM  04/11/2017 3:19:58 PM
1             1.80          1
2            106203690         1  12/15/2017 7:26:56 AM  12/15/2017 7:34:08 AM
1             1.00          1
3            38942136          2  05/07/2017 1:17:59 PM  05/07/2017 1:48:14 PM
1             3.70          1
4            30841670          2  04/15/2017 11:32:20 PM  04/15/2017 11:49:03 PM
1             4.37          1
5            23345809          2  03/25/2017 8:34:11 PM  03/25/2017 8:42:11 PM
6             2.30          1
6            37660487          2  05/03/2017 7:04:09 PM  05/03/2017 8:03:47 PM
```

1		12.83		1		
7	69059411		2	08/15/2017 5:41:06 PM	08/15/2017 6:03:05 PM	
1		2.98		1		
8	8433159		2	02/04/2017 4:17:07 PM	02/04/2017 4:29:14 PM	
1		1.20		1		
9	95294817		1	11/10/2017 3:20:29 PM	11/10/2017 3:40:55 PM	
1		1.60		1		
store_and_fwd_flag PULocationID DOLocationID payment_type fare_amount						
extra mta_tax tip_amount tolls_amount \						
0		N	100	231	1	13.0
0.0	0.5	2.76	0.0			
1		N	186	43	1	16.0
0.0	0.5	4.00	0.0			
2		N	262	236	1	6.5
0.0	0.5	1.45	0.0			
3		N	188	97	1	20.5
0.0	0.5	6.39	0.0			
4		N	4	112	2	16.5
0.5	0.5	0.00	0.0			
5		N	161	236	1	9.0
0.5	0.5	2.06	0.0			
6		N	79	241	1	47.5
1.0	0.5	9.86	0.0			
7		N	237	114	1	16.0
1.0	0.5	1.78	0.0			
8		N	234	249	2	9.0
0.0	0.5	0.00	0.0			
9		N	239	237	1	13.0
0.0	0.5	2.75	0.0			
improvement_surcharge total_amount						
0		0.3	16.56			
1		0.3	20.80			
2		0.3	8.75			
3		0.3	27.69			
4		0.3	17.80			
5		0.3	12.36			
6		0.3	59.16			
7		0.3	19.58			
8		0.3	9.80			
9		0.3	16.55			

Inspect the first few rows of `nyc_preds_means`.

```
[126]: # Inspect the first few rows of `nyc_preds_means`  
nyc_preds_means.head()
```

```
[126]:   mean_duration  mean_distance  predicted_fare
0        22.847222      3.521667    16.434245
1        24.470370      3.108889    16.052218
2         7.250000      0.881429     7.053706
3       30.250000      3.700000    18.731650
4       14.616667      4.435000    15.845642
```

Join the two dataframes Join the two dataframes using a method of your choice.

```
[164]: df0 = df0.merge(nyc_preds_means,
                     left_index=True,
                     right_index=True)

df0.head()
```

```
[164]:   Unnamed: 0  VendorID      tpep_pickup_datetime      tpep_dropoff_datetime
passenger_count  trip_distance  RatecodeID \
0            24870114          2  03/25/2017 8:55:43 AM  03/25/2017 9:09:47 AM
6             3.34           1
1            35634249          1  04/11/2017 2:53:28 PM  04/11/2017 3:19:58 PM
1             1.80           1
2            106203690         1  12/15/2017 7:26:56 AM  12/15/2017 7:34:08 AM
1             1.00           1
3            38942136          2  05/07/2017 1:17:59 PM  05/07/2017 1:48:14 PM
1             3.70           1
4            30841670          2  04/15/2017 11:32:20 PM  04/15/2017 11:49:03 PM
1             4.37           1

      store_and_fwd_flag  PULocationID  DOLocationID  payment_type  fare_amount
extra  mta_tax  tip_amount  tolls_amount \
0                 N          100          231              1        13.0
0.0     0.5      2.76          0.0
1                 N          186           43              1        16.0
0.0     0.5      4.00          0.0
2                 N          262          236              1        6.5
0.0     0.5      1.45          0.0
3                 N          188           97              1        20.5
0.0     0.5      6.39          0.0
4                 N              4          112              2        16.5
0.5     0.5      0.00          0.0

      improvement_surcharge  total_amount  mean_duration  mean_distance
predicted_fare
0                  0.3        16.56    22.847222      3.521667
1                  0.3        20.80    24.470370      3.108889
16.434245
16.052218
```

2	0.3	8.75	7.250000	0.881429
7.053706				
3	0.3	27.69	30.250000	3.700000
18.731650				
4	0.3	17.80	14.616667	4.435000
15.845642				

4.2 PACE: Analyze

Consider the questions in your PACE Strategy Document to reflect on the Analyze stage.

4.2.1 Task 2. Feature engineering

You have already prepared much of this data and performed exploratory data analysis (EDA) in previous courses.

Call `info()` on the new combined dataframe.

[165]: `df0.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        22699 non-null   int64  
 1   VendorID          22699 non-null   int64  
 2   tpep_pickup_datetime  22699 non-null   object  
 3   tpep_dropoff_datetime 22699 non-null   object  
 4   passenger_count    22699 non-null   int64  
 5   trip_distance      22699 non-null   float64 
 6   RatecodeID         22699 non-null   int64  
 7   store_and_fwd_flag 22699 non-null   object  
 8   PULocationID      22699 non-null   int64  
 9   DOLocationID       22699 non-null   int64  
 10  payment_type       22699 non-null   int64  
 11  fare_amount        22699 non-null   float64 
 12  extra              22699 non-null   float64 
 13  mta_tax             22699 non-null   float64 
 14  tip_amount          22699 non-null   float64 
 15  tolls_amount        22699 non-null   float64 
 16  improvement_surcharge 22699 non-null   float64 
 17  total_amount        22699 non-null   float64 
 18  mean_duration       22699 non-null   float64 
 19  mean_distance        22699 non-null   float64 
 20  predicted_fare      22699 non-null   float64 
dtypes: float64(11), int64(7), object(3)
memory usage: 3.6+ MB
```

You know from your EDA that customers who pay cash generally have a tip amount of \$0. To meet the modeling objective, you'll need to sample the data to select only the customers who pay with credit card.

Copy `df0` and assign the result to a variable called `df1`. Then, use a Boolean mask to filter `df1` so it contains only customers who paid with credit card.

```
[166]: df1 = df0[df0['payment_type']==1]
```

Target Notice that there isn't a column that indicates tip percent, which is what you need to create the target variable. You'll have to engineer it.

Add a `tip_percent` column to the dataframe by performing the following calculation:

$$\text{tip percent} = \frac{\text{tip amount}}{\text{total amount} - \text{tip amount}}$$

Round the result to three places beyond the decimal. **This is an important step.** It affects how many customers are labeled as generous tippers. In fact, without performing this step, approximately 1,800 people who do tip 20% would be labeled as not generous.

To understand why, you must consider how floats work. Computers make their calculations using floating-point arithmetic (hence the word “float”). Floating-point arithmetic is a system that allows computers to express both very large numbers and very small numbers with a high degree of precision, encoded in binary. However, precision is limited by the number of bits used to represent a number, which is generally 32 or 64, depending on the capabilities of your operating system.

This comes with limitations in that sometimes calculations that should result in clean, precise values end up being encoded as very long decimals. Take, for example, the following calculation:

```
[127]: # Run this cell  
1.1 + 2.2
```

```
[127]: 3.3000000000000003
```

Notice the three that is 16 places to the right of the decimal. As a consequence, if you were to then have a step in your code that identifies values 3.3, this would not be included in the result. Therefore, whenever you perform a calculation to compute a number that is then used to make an important decision or filtration, round the number. How many degrees of precision you round to is your decision, which should be based on your use case.

Refer to this [guide](#) for more information related to floating-point arithmetic.

```
[167]: df1['tip_percent'] = round(df1['tip_amount'] / (df1['total_amount'] -  
        ↴df1['tip_amount']), 3)
```

Now create another column called `generous`. This will be the target variable. The column should be a binary indicator of whether or not a customer tipped 20% (0=no, 1=yes).

1. Begin by making the `generous` column a copy of the `tip_percent` column.

2. Reassign the column by converting it to Boolean (True/False).
3. Reassign the column by converting Boolean to binary (1/0).

```
[168]: # Create 'generous' col (target)
df1['generous'] = df1['tip_percent']
df1['generous'] = (df1['generous'] >= 0.2)
df1['generous'] = df1['generous'].astype(int)
```

HINT

To convert from Boolean to binary, use `.astype(int)` on the column.

[78]:

Create day column Next, you're going to be working with the pickup and dropoff columns.

Convert the `tpep_pickup_datetime` and `tpep_dropoff_datetime` columns to datetime.

```
[169]: # Convert pickup and dropoff cols to datetime
df1['tpep_pickup_datetime'] = pd.to_datetime(df1['tpep_pickup_datetime'],
                                             format='%m/%d/%Y %I:%M:%S %p')
df1['tpep_dropoff_datetime'] = pd.to_datetime(df1['tpep_dropoff_datetime'],
                                             format='%m/%d/%Y %I:%M:%S %p')
```

Create a `day` column that contains only the day of the week when each passenger was picked up. Then, convert the values to lowercase.

```
[170]: # Create a 'day' col
df1['day'] = df1['tpep_pickup_datetime'].dt.day_name().str.lower()
df1[['tpep_pickup_datetime', 'day']].head()
```

```
[170]: tpep_pickup_datetime      day
0  2017-03-25 08:55:43  saturday
1  2017-04-11 14:53:28  tuesday
2  2017-12-15 07:26:56  friday
3  2017-05-07 13:17:59  sunday
5  2017-03-25 20:34:11  saturday
```

HINT

To convert to day name, use `dt.day_name()` on the column.

Create time of day columns Next, engineer four new columns that represent time of day bins. Each column should contain binary values (0=no, 1=yes) that indicate whether a trip began (picked up) during the following times:

```
am_rush = [06:00–10:00)
daytime = [10:00–16:00)
pm_rush = [16:00–20:00)
nighttime = [20:00–06:00)
```

To do this, first create the four columns. For now, each new column should be identical and contain the same information: the hour (only) from the tpep_pickup_datetime column.

```
[171]: # Create 'am_rush' col
df1['am_rush'] = df1['tpep_pickup_datetime'].dt.hour

# Create 'daytime' col
df1['daytime'] = df1['tpep_pickup_datetime'].dt.hour

# Create 'pm_rush' col
df1['pm_rush'] = df1['tpep_pickup_datetime'].dt.hour

# Create 'nighttime' col
df1['nighttime'] = df1['tpep_pickup_datetime'].dt.hour
```

You'll need to write four functions to convert each new column to binary (0/1). Begin with am_rush. Complete the function so if the hour is between [06:00–10:00), it returns 1, otherwise, it returns 0.

```
[172]: # Define 'am_rush()' conversion function [06:00-10:00)
def am_rush(hour):
    if 6 <= hour['am_rush'] < 10:
        val = 1
    else:
        val = 0
    return val
```

Now, apply the am_rush() function to the am_rush series to perform the conversion. Print the first five values of the column to make sure it did what you expected it to do.

Note: Be careful! If you run this cell twice, the function will be reapplied and the values will all be changed to 0.

```
[176]: # Apply 'am_rush' function to the 'am_rush' series
df1['am_rush'] = df1.apply(am_rush, axis=1)
df1['am_rush'].head()
```

```
[176]: 0      0
       1      0
       2      0
       3      0
       5      0
Name: am_rush, dtype: int64
```

Write functions to convert the three remaining columns and apply them to their respective series.

```
[177]: # Define 'daytime()' conversion function [10:00-16:00)
def daytime(hour):
    if 10 <= hour['daytime'] < 16:
        val = 1
```

```
    else:  
        val = 0  
    return val
```

```
[178]: # Apply 'daytime' function to the 'daytime' series  
df1['daytime'] = df1.apply(daytime, axis=1)
```

```
[179]: # Define 'pm_rush()' conversion function [16:00-20:00)  
def pm_rush(hour):  
    if 16 <= hour['pm_rush'] < 20:  
        val = 1  
    else:  
        val = 0  
    return val
```

```
[180]: # Apply 'pm_rush' function to the 'pm_rush' series  
df1['pm_rush'] = df1.apply(pm_rush, axis=1)
```

```
[181]: # Define 'nighttime()' conversion function [20:00-06:00)  
def nighttime(hour):  
    if 20 <= hour['nighttime'] < 24:  
        val = 1  
    elif 0 <= hour['nighttime'] < 6:  
        val = 1  
    else:  
        val = 0  
    return val
```

```
[182]: # Apply 'nighttime' function to the 'nighttime' series  
df1['nighttime'] = df1.apply(nighttime, axis=1)
```

Create month column Now, create a `month` column that contains only the abbreviated name of the month when each passenger was picked up, then convert the result to lowercase.

HINT

Refer to the [strftime cheatsheet](#) for help.

```
[183]: # Create 'month' col  
df1['month'] = df1['tpep_pickup_datetime'].dt.strftime('%b').str.lower()
```

Examine the first five rows of your dataframe.

```
[184]: df1.head()
```

```
[184]: Unnamed: 0  VendorID tpep_pickup_datetime tpep_dropoff_datetime  
passenger_count  trip_distance RatecodeID  \  
0      24870114          2 2017-03-25 08:55:43  2017-03-25 09:09:47
```

6	3.34	1				
1	35634249	1	2017-04-11 14:53:28	2017-04-11 15:19:58		
1	1.80	1				
2	106203690	1	2017-12-15 07:26:56	2017-12-15 07:34:08		
1	1.00	1				
3	38942136	2	2017-05-07 13:17:59	2017-05-07 13:48:14		
1	3.70	1				
5	23345809	2	2017-03-25 20:34:11	2017-03-25 20:42:11		
6	2.30	1				
store_and_fwd_flag PUlocationID DOLocationID payment_type fare_amount						
extra mta_tax tip_amount tolls_amount \						
0	N	100	231	1	13.0	
0.0	0.5	2.76	0.0			
1	N	186	43	1	16.0	
0.0	0.5	4.00	0.0			
2	N	262	236	1	6.5	
0.0	0.5	1.45	0.0			
3	N	188	97	1	20.5	
0.0	0.5	6.39	0.0			
5	N	161	236	1	9.0	
0.5	0.5	2.06	0.0			
improvement_surcharge total_amount mean_duration mean_distance						
predicted_fare tip_percent generous day \						
0	0.3	16.56	22.847222	3.521667		
16.434245	0.200	1 saturday				
1	0.3	20.80	24.470370	3.108889		
16.052218	0.238	1 tuesday				
2	0.3	8.75	7.250000	0.881429		
7.053706	0.199	0 friday				
3	0.3	27.69	30.250000	3.700000		
18.731650	0.300	1 sunday				
5	0.3	12.36	11.855376	2.052258		
10.441351	0.200	1 saturday				
am_rush daytime pm_rush nighttime month						
0	0	0	0	0 mar		
1	0	0	0	0 apr		
2	0	0	0	0 dec		
3	0	0	0	0 may		
5	0	0	0	1 mar		

Drop columns Drop redundant and irrelevant columns as well as those that would not be available when the model is deployed. This includes information like payment type, trip distance, tip amount, tip percentage, total amount, toll amount, etc. The target variable (`generous`) must remain in the data because it will get isolated as the y data for modeling.

```
[185]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 15265 entries, 0 to 22698
Data columns (total 29 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        15265 non-null   int64  
 1   VendorID         15265 non-null   int64  
 2   tpep_pickup_datetime  15265 non-null   datetime64[ns]
 3   tpep_dropoff_datetime 15265 non-null   datetime64[ns]
 4   passenger_count    15265 non-null   int64  
 5   trip_distance      15265 non-null   float64 
 6   RatecodeID         15265 non-null   int64  
 7   store_and_fwd_flag 15265 non-null   object  
 8   PULocationID       15265 non-null   int64  
 9   DOLocationID       15265 non-null   int64  
 10  payment_type       15265 non-null   int64  
 11  fare_amount         15265 non-null   float64 
 12  extra              15265 non-null   float64 
 13  mta_tax             15265 non-null   float64 
 14  tip_amount          15265 non-null   float64 
 15  tolls_amount        15265 non-null   float64 
 16  improvement_surcharge 15265 non-null   float64 
 17  total_amount        15265 non-null   float64 
 18  mean_duration       15265 non-null   float64 
 19  mean_distance        15265 non-null   float64 
 20  predicted_fare      15265 non-null   float64 
 21  tip_percent          15262 non-null   float64 
 22  generous            15265 non-null   int64  
 23  day                 15265 non-null   object  
 24  am_rush              15265 non-null   int64  
 25  daytime              15265 non-null   int64  
 26  pm_rush              15265 non-null   int64  
 27  nighttime            15265 non-null   int64  
 28  month                15265 non-null   object  
dtypes: datetime64[ns](2), float64(12), int64(12), object(3)
memory usage: 3.5+ MB
```

```
[186]: # Drop columns
```

```
drop_cols = ['Unnamed: 0', 'tpep_pickup_datetime', 'tpep_dropoff_datetime',
             'payment_type', 'trip_distance', 'store_and_fwd_flag', □
             ↵'payment_type',
             'fare_amount', 'extra', 'mta_tax', 'tip_amount', 'tolls_amount',
             'improvement_surcharge', 'total_amount', 'tip_percent']

df1 = df1.drop(drop_cols, axis=1)
```

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 15265 entries, 0 to 22698
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   VendorID        15265 non-null   int64  
 1   passenger_count 15265 non-null   int64  
 2   RatecodeID      15265 non-null   int64  
 3   PULocationID   15265 non-null   int64  
 4   DOLocationID   15265 non-null   int64  
 5   mean_duration   15265 non-null   float64 
 6   mean_distance   15265 non-null   float64 
 7   predicted_fare  15265 non-null   float64 
 8   generous        15265 non-null   int64  
 9   day              15265 non-null   object  
 10  am_rush         15265 non-null   int64  
 11  daytime         15265 non-null   int64  
 12  pm_rush         15265 non-null   int64  
 13  nighttime        15265 non-null   int64  
 14  month            15265 non-null   object  
dtypes: float64(3), int64(10), object(2)
memory usage: 1.9+ MB
```

Variable encoding Many of the columns are categorical and will need to be dummmied (converted to binary). Some of these columns are numeric, but they actually encode categorical information, such as `RatecodeID` and the pickup and dropoff locations. To make these columns recognizable to the `get_dummies()` function as categorical variables, you'll first need to convert them to `type(str)`.

1. Define a variable called `cols_to_str`, which is a list of the numeric columns that contain categorical information and must be converted to string: `RatecodeID`, `PULocationID`, `DOLocationID`.
2. Write a for loop that converts each column in `cols_to_str` to string.

```
[187]: # 1. Define list of cols to convert to string
cols_to_str = ['RatecodeID', 'PULocationID', 'DOLocationID', 'VendorID']

# 2. Convert each column to string
for col in cols_to_str:
    df1[col] = df1[col].astype('str')
```

```
[150]: df1[cols_to_str].dtypes
```

```
RatecodeID      object
PULocationID   object
DOLocationID   object
dtype: object
```

HINT

To convert to string, use `astype(str)` on the column.

Now convert all the categorical columns to binary.

1. Call `get_dummies()` on the dataframe and assign the results back to a new dataframe called `df2`.

```
[190]: # Convert categoricals to binary
df2 = pd.get_dummies(df1, drop_first=True)
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 15265 entries, 0 to 22698
Columns: 347 entries, passenger_count to month_sep
dtypes: bool(338), float64(3), int64(6)
memory usage: 6.1 MB
```

```
[191]: df2.info()
df2.head()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 15265 entries, 0 to 22698
Columns: 347 entries, passenger_count to month_sep
dtypes: bool(338), float64(3), int64(6)
memory usage: 6.1 MB
```

```
[191]:   passenger_count  mean_duration  mean_distance  predicted_fare  generous
am_rush      daytime      pm_rush    nighttme \
0            6        22.847222      3.521667     16.434245      1
0            0            0            0
1            1        24.470370      3.108889     16.052218      1
0            0            0            0
2            1        7.250000      0.881429      7.053706      0
0            0            0            0
3            1        30.250000      3.700000     18.731650      1
0            0            0            0
5            6        11.855376      2.052258     10.441351      1
0            0            0            1

   VendorID_2  RatecodeID_2  RatecodeID_3  RatecodeID_4  RatecodeID_5
RatecodeID_99  PULocationID_10 \
0       True      False      False      False      False
False      False
1      False      False      False      False      False
False      False
2      False      False      False      False      False
False      False
```

3	True	False	False	False	False
False		False			
5	True	False	False	False	False
False		False			
	PULocationID_100	PULocationID_106	PULocationID_107	PULocationID_112	
	PULocationID_113	PULocationID_114	\		
0	True	False		False	False
False		False			
1	False	False		False	False
False		False			
2	False	False		False	False
False		False			
3	False	False		False	False
False		False			
5	False	False		False	False
False		False			
	PULocationID_116	PULocationID_12	PULocationID_123	PULocationID_125	
	PULocationID_127	PULocationID_128	\		
0	False	False		False	False
False		False			
1	False	False		False	False
False		False			
2	False	False		False	False
False		False			
3	False	False		False	False
False		False			
5	False	False		False	False
False		False			
	PULocationID_129	PULocationID_13	PULocationID_130	PULocationID_131	
	PULocationID_132	PULocationID_133	\		
0	False	False		False	False
False		False			
1	False	False		False	False
False		False			
2	False	False		False	False
False		False			
3	False	False		False	False
False		False			
5	False	False		False	False
False		False			
	PULocationID_134	PULocationID_135	PULocationID_137	PULocationID_138	
	PULocationID_140	PULocationID_141	\		
0	False	False		False	False

False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False

	PULocationID_142	PULocationID_143	PULocationID_144	PULocationID_145
PULocationID_146	PULocationID_148	\		
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False

	PULocationID_151	PULocationID_152	PULocationID_153	PULocationID_158
PULocationID_161	PULocationID_162	\		
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
True	False	False	False	False

	PULocationID_163	PULocationID_164	PULocationID_166	PULocationID_17
PULocationID_170	PULocationID_173	\		
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False

False	False			
PULocationID_179	PULocationID_181	PULocationID_186	PULocationID_188	
PULocationID_189	PULocationID_190	\		
0	False	False	False	False
False	False			
1	False	False	True	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	True
False	False			
5	False	False	False	False
False	False			
PULocationID_193	PULocationID_196	PULocationID_208	PULocationID_209	
PULocationID_211	PULocationID_213	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			
PULocationID_216	PULocationID_218	PULocationID_223	PULocationID_224	
PULocationID_225	PULocationID_226	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			
PULocationID_229	PULocationID_230	PULocationID_231	PULocationID_232	
PULocationID_233	PULocationID_234	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			

2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			

PULocationID_236	PULocationID_237	PULocationID_238	PULocationID_239	
PULocationID_24	PULocationID_243	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			

PULocationID_244	PULocationID_246	PULocationID_247	PULocationID_249	
PULocationID_25	PULocationID_255	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			

PULocationID_256	PULocationID_258	PULocationID_260	PULocationID_261	
PULocationID_262	PULocationID_263	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
True	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			

PULocationID_264 PULocationID_265 PULocationID_28 PULocationID_33

PULocationID_35	PULocationID_36	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			
PULocationID_37	PULocationID_4	PULocationID_40	PULocationID_41	
PULocationID_42	PULocationID_43	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			
PULocationID_45	PULocationID_48	PULocationID_49	PULocationID_50	
PULocationID_52	PULocationID_57	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			
PULocationID_61	PULocationID_62	PULocationID_65	PULocationID_66	
PULocationID_68	PULocationID_7	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False

False	False			
5	False	False	False	False
False	False			
	PULocationID_70	PULocationID_74	PULocationID_75	PULocationID_79
	PULocationID_80	PULocationID_82	\	
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			
	PULocationID_87	PULocationID_88	PULocationID_90	PULocationID_91
	PULocationID_92	PULocationID_93	\	
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			
	PULocationID_95	PULocationID_97	DOLocationID_10	DOLocationID_100
	DOLocationID_102	DOLocationID_106	\	
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			
	DOLocationID_107	DOLocationID_11	DOLocationID_112	DOLocationID_113
	DOLocationID_114	DOLocationID_116	\	
0	False	False	False	False
False	False			

1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False

DOLocationID_117	DOLocationID_118	DOLocationID_119	DOLocationID_12	
DOLocationID_120	DOLocationID_121	\		
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False

DOLocationID_123	DOLocationID_124	DOLocationID_125	DOLocationID_126	
DOLocationID_127	DOLocationID_129	\		
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False

DOLocationID_13	DOLocationID_130	DOLocationID_131	DOLocationID_132	
DOLocationID_133	DOLocationID_134	\		
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False

DOLocationID_135 DOLocationID_136 DOLocationID_137 DOLocationID_138
DOLocationID_14 DOLocationID_140 \

0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False

DOLocationID_141 DOLocationID_142 DOLocationID_143 DOLocationID_144
DOLocationID_145 DOLocationID_146 \

0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False

DOLocationID_147 DOLocationID_148 DOLocationID_15 DOLocationID_151
DOLocationID_152 DOLocationID_153 \

0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False

DOLocationID_157 DOLocationID_158 DOLocationID_159 DOLocationID_16
DOLocationID_160 DOLocationID_161 \

0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False

False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_162	DOLocationID_163	DOLocationID_164	DOLocationID_166
	DOLocationID_168	DOLocationID_169	DOLocationID_164	DOLocationID_166
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_17	DOLocationID_170	DOLocationID_173	DOLocationID_174
	DOLocationID_175	DOLocationID_177	DOLocationID_173	DOLocationID_174
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_178	DOLocationID_179	DOLocationID_180	DOLocationID_181
	DOLocationID_182	DOLocationID_183	DOLocationID_180	DOLocationID_181
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_186	DOLocationID_188	DOLocationID_189	DOLocationID_19
	DOLocationID_192	DOLocationID_193	DOLocationID_189	DOLocationID_19

0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_194	DOLocationID_195	DOLocationID_196	DOLocationID_197
	DOLocationID_198	DOLocationID_200	\	
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_202	DOLocationID_208	DOLocationID_209	DOLocationID_21
	DOLocationID_210	DOLocationID_211	\	
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_212	DOLocationID_213	DOLocationID_216	DOLocationID_217
	DOLocationID_218	DOLocationID_22	\	
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False

5	False	False	False	False
False	False			
	DOLocationID_220	DOLocationID_223	DOLocationID_224	DOLocationID_225
	DOLocationID_226	DOLocationID_228 \		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			
	DOLocationID_229	DOLocationID_23	DOLocationID_230	DOLocationID_231
	DOLocationID_232	DOLocationID_233 \		
0	False	False	False	True
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			
	DOLocationID_234	DOLocationID_235	DOLocationID_236	DOLocationID_237
	DOLocationID_238	DOLocationID_239 \		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	True	False
False	False			
3	False	False	False	False
False	False			
5	False	False	True	False
False	False			
	DOLocationID_24	DOLocationID_240	DOLocationID_241	DOLocationID_242
	DOLocationID_243	DOLocationID_244 \		
0	False	False	False	False
False	False			
1	False	False	False	False

False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_246	DOLocationID_247	DOLocationID_248	DOLocationID_249
	DOLocationID_25	DOLocationID_252	\	
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_255	DOLocationID_256	DOLocationID_257	DOLocationID_259
	DOLocationID_26	DOLocationID_260	\	
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_261	DOLocationID_262	DOLocationID_263	DOLocationID_264
	DOLocationID_265	DOLocationID_28	\	
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False

DOLocationID_29	DOLocationID_32	DOLocationID_33	DOLocationID_36	
DOLocationID_37	DOLocationID_39	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			

DOLocationID_4	DOLocationID_40	DOLocationID_41	DOLocationID_42	
DOLocationID_43	DOLocationID_45	\		
0	False	False	False	False
False	False			
1	False	False	False	False
True	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			

DOLocationID_47	DOLocationID_48	DOLocationID_49	DOLocationID_50	
DOLocationID_51	DOLocationID_52	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			
3	False	False	False	False
False	False			
5	False	False	False	False
False	False			

DOLocationID_53	DOLocationID_54	DOLocationID_55	DOLocationID_56	
DOLocationID_61	DOLocationID_62	\		
0	False	False	False	False
False	False			
1	False	False	False	False
False	False			
2	False	False	False	False
False	False			

3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_63	DOLocationID_64	DOLocationID_65	DOLocationID_66
	DOLocationID_67	DOLocationID_68	\	
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_69	DOLocationID_7	DOLocationID_70	DOLocationID_71
	DOLocationID_72	DOLocationID_74	\	
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_75	DOLocationID_76	DOLocationID_77	DOLocationID_79
	DOLocationID_80	DOLocationID_81	\	
0	False	False	False	False
False	False	False	False	False
1	False	False	False	False
False	False	False	False	False
2	False	False	False	False
False	False	False	False	False
3	False	False	False	False
False	False	False	False	False
5	False	False	False	False
False	False	False	False	False
	DOLocationID_82	DOLocationID_83	DOLocationID_85	DOLocationID_86
	DOLocationID_87	DOLocationID_88	\	
0	False	False	False	False

False	False	False	False	False	False		
1	False	False	False	False	False		
False	False	False	False	False	False		
2	False	False	False	False	False		
False	False	False	False	False	False		
3	False	False	False	False	False		
False	False	False	False	False	False		
5	False	False	False	False	False		
False	False	False	False	False	False		
	DOLocationID_89	DOLocationID_9	DOLocationID_90	DOLocationID_91			
	DOLocationID_92	DOLocationID_93	\				
0	False	False	False	False	False		
False	False	False	False	False	False		
1	False	False	False	False	False		
False	False	False	False	False	False		
2	False	False	False	False	False		
False	False	False	False	False	False		
3	False	False	False	False	False		
False	False	False	False	False	False		
5	False	False	False	False	False		
False	False	False	False	False	False		
	DOLocationID_95	DOLocationID_97	day_monday	day_saturday	day_sunday		
	day_thursday	day_tuesday	day_wednesday	\			
0	False	False	False	False	True		
False	False	False	False	False	False		
1	False	False	False	False	False		
False	True	False	False	False	False		
2	False	False	False	False	False		
False	False	False	False	False	False		
3	False	True	False	False	False		
False	False	False	False	False	True		
5	False	False	False	False	False		
False	False	False	False	True	False		
	month_aug	month_dec	month_feb	month_jan	month_jul	month_jun	month_mar
	month_may	month_nov	month_oct	\			
0	False	False	False	False	False	False	True
False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False
False	False	False	False	False	False	False	False
2	False	True	False	False	False	False	False
False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False
True	False	False	False	False	False	False	False
5	False	False	False	False	False	False	True

```
False      False      False  
month_sep  
0    False  
1    False  
2    False  
3    False  
5    False
```

```
[193]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 15265 entries, 0 to 22698  
Columns: 347 entries, passenger_count to month_sep  
dtypes: bool(338), float64(3), int64(6)  
memory usage: 6.1 MB
```

Evaluation metric Before modeling, you must decide on an evaluation metric.

1. Examine the class balance of your target variable.

```
[189]: # Get class balance of 'generous' col  
df2['generous'].value_counts(normalize=True)
```

```
[189]: generous  
1    0.526368  
0    0.473632  
Name: proportion, dtype: float64
```

A little over half of the customers in this dataset were “generous” (tipped 20%). The dataset is very nearly balanced.

To determine a metric, consider the cost of both kinds of model error:

- * False positives (the model predicts a tip 20%, but the customer does not give one)
- * False negatives (the model predicts a tip < 20%, but the customer gives more)

False positives are worse for cab drivers, because they would pick up a customer expecting a good tip and then not receive one, frustrating the driver.

False negatives are worse for customers, because a cab driver would likely pick up a different customer who was predicted to tip more—even when the original customer would have tipped generously.

The stakes are relatively even. You want to help taxi drivers make more money, but you don’t want this to anger customers. Your metric should weigh both precision and recall equally. Which metric is this?

Since the dataset is nearly balanced and both false positives and false negatives carry similar costs, the best evaluation metric is the F1 score, which balances precision and recall equally.

4.3 PACE: Construct

Consider the questions in your PACE Strategy Document to reflect on the Construct stage.

4.3.1 Task 3. Modeling

Split the data Now you're ready to model. The only remaining step is to split the data into features/target variable and training/testing data.

1. Define a variable `y` that isolates the target variable (`generous`).
2. Define a variable `X` that isolates the features.
3. Split the data into training and testing sets. Put 20% of the samples into the test set, stratify the data, and set the random state.

```
[194]: # Isolate target variable (y)
y = df2['generous']

# Isolate the features (X)
X = df2.drop('generous', axis=1)

# Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, u
↪test_size=0.2, random_state=42)
```

Random forest Begin with using `GridSearchCV` to tune a random forest model.

1. Instantiate the random forest classifier `rf` and set the random state.
2. Create a dictionary `cv_params` of any of the following hyperparameters and their corresponding values to tune. The more you tune, the better your model will fit the data, but the longer it will take.
 - `max_depth`
 - `max_features`
 - `max_samples`
 - `min_samples_leaf`
 - `min_samples_split`
 - `n_estimators`
3. Define a set `scoring` of scoring metrics for `GridSearch` to capture (precision, recall, F1 score, and accuracy).
4. Instantiate the `GridSearchCV` object `rf1`. Pass to it as arguments:
 - `estimator=rf`
 - `param_grid=cv_params`
 - `scoring=scoring`
 - `cv`: define the number of cross-validation folds you want (`cv=_`)
 - `refit`: indicate which evaluation metric you want to use to select the model (`refit=_`)

Note: `refit` should be set to '`f1`'.

```
[195]: # 1. Instantiate the random forest classifier
rf = RandomForestClassifier(random_state=42)

# 2. Create a dictionary of hyperparameters to tune
# Note that this example only contains 1 value for each parameter for
# simplicity,
# but you should assign a dictionary with ranges of values
cv_params = {'max_depth': [None],
             'max_features': [1.0],
             'max_samples': [0.7],
             'min_samples_leaf': [1],
             'min_samples_split': [2],
             'n_estimators': [300]
            }

# 3. Define a list of scoring metrics to capture
scoring = ['accuracy', 'precision', 'recall', 'f1']

# 4. Instantiate the GridSearchCV object
rf1 = GridSearchCV(rf, cv_params, scoring=scoring, cv=4, refit='f1')
```

```
[158]:
```

Now fit the model to the training data. Note that, depending on how many options you include in your search grid and the number of cross-validation folds you select, this could take a very long time—even hours. If you use 4-fold validation and include only one possible value for each hyperparameter and grow 300 trees to full depth, it should take about 5 minutes. If you add another value for GridSearch to check for, say, `min_samples_split` (so all hyperparameters now have 1 value except for `min_samples_split`, which has 2 possibilities), it would double the time to ~10 minutes. Each additional parameter would approximately double the time.

```
[196]: %%time
rf1.fit(X_train, y_train)
```

```
CPU times: user 4min 40s, sys: 247 ms, total: 4min 40s
Wall time: 4min 40s
```

```
[196]: GridSearchCV(cv=4, estimator=RandomForestClassifier(random_state=42),
                     param_grid={'max_depth': [None], 'max_features': [1.0],
                                 'max_samples': [0.7], 'min_samples_leaf': [1],
                                 'min_samples_split': [2], 'n_estimators': [300]},
                     refit='f1', scoring=['accuracy', 'precision', 'recall', 'f1'])
```

HINT

If you get a warning that a metric is 0 due to no predicted samples, think about how many features you're sampling with `max_features`. How many features are in the dataset? How many are likely

predictive enough to give good predictions within the number of splits you've allowed (determined by the `max_depth` hyperparameter)? Consider increasing `max_features`.

If you want, use `pickle` to save your models and read them back in. This can be particularly helpful when performing a search over many possible hyperparameter values.

```
[202]: import pickle

# Define a path to the folder where you want to save the model
path = '/home/jovyan/work/'
```

```
[199]: def write_pickle(path, model_object, save_name:str):
    """
    save_name is a string.
    """
    with open(path + save_name + '.pickle', 'wb') as to_write:
        pickle.dump(model_object, to_write)
```

```
[203]: def read_pickle(path, saved_model_name:str):
    """
    saved_model_name is a string.
    """
    with open(path + saved_model_name + '.pickle', 'rb') as to_read:
        model = pickle.load(to_read)

    return model
```

Examine the best average score across all the validation folds.

```
[204]: # Examine best score
rf1.best_score_
```

```
[204]: 0.7136183532391456
```

Examine the best combination of hyperparameters.

```
[205]: rf1.best_params_
```

```
[205]: {'max_depth': None,
         'max_features': 1.0,
         'max_samples': 0.7,
         'min_samples_leaf': 1,
         'min_samples_split': 2,
         'n_estimators': 300}
```

Use the `make_results()` function to output all of the scores of your model. Note that it accepts three arguments.

HINT

To learn more about how this function accesses the cross-validation results, refer to the [GridSearchCV scikit-learn documentation](#) for the `cv_results_` attribute.

```
[206]: def make_results(model_name:str, model_object, metric:str):
    """
    Arguments:
        model_name (string): what you want the model to be called in the output
        ↪table
        model_object: a fit GridSearchCV object
        metric (string): precision, recall, f1, or accuracy

    Returns a pandas df with the F1, recall, precision, and accuracy scores
    for the model with the best mean 'metric' score across all validation folds.
    """

    # Create dictionary that maps input metric to actual metric name in
    ↪GridSearchCV
    metric_dict = {'precision': 'mean_test_precision',
                   'recall': 'mean_test_recall',
                   'f1': 'mean_test_f1',
                   'accuracy': 'mean_test_accuracy',
                   }

    # Get all the results from the CV and put them in a df
    cv_results = pd.DataFrame(model_object.cv_results_)

    # Isolate the row of the df with the max(metric) score
    best_estimator_results = cv_results.iloc[cv_results[metric_dict[metric]].idxmax(), :]

    # Extract Accuracy, precision, recall, and f1 score from that row
    f1 = best_estimator_results.mean_test_f1
    recall = best_estimator_results.mean_test_recall
    precision = best_estimator_results.mean_test_precision
    accuracy = best_estimator_results.mean_test_accuracy

    # Create table of results
    table = pd.DataFrame({'model': [model_name],
                          'precision': [precision],
                          'recall': [recall],
                          'F1': [f1],
                          'accuracy': [accuracy],
                          },
                          )

    return table
```

Call `make_results()` on the `GridSearch` object.

```
[207]: results = make_results('RF CV', rf1, 'f1')
results
```

```
[207]:    model  precision      recall          F1   accuracy
0  RF CV      0.67476  0.757467  0.713618  0.680151
```

Your results should produce an acceptable model across the board. Typically scores of 0.65 or better are considered acceptable, but this is always dependent on your use case. Optional: try to improve the scores. It's worth trying, especially to practice searching over different hyperparameters.

HINT

For example, if the available values for `min_samples_split` were [2, 3, 4] and GridSearch identified the best value as 4, consider trying [4, 5, 6] this time.

Use your model to predict on the test data. Assign the results to a variable called `rf_preds`.

HINT

You cannot call `predict()` on the `GridSearchCV` object directly. You must call it on the `best_estimator_`.

For this project, you will use several models to predict on the test data. Remember that this decision comes with a trade-off. What is the benefit of this? What is the drawback?

==> ENTER YOUR RESPONSE HERE

```
[208]: # Get scores on test data
rf_preds = rf1.best_estimator_.predict(X_test)
```

Use the below `get_test_scores()` function you will use to output the scores of the model on the test data.

```
[209]: def get_test_scores(model_name:str, preds, y_test_data):
    """
    Generate a table of test scores.

    In:
    model_name (string): Your choice: how the model will be named in the output
    ↴table
        preds: numpy array of test predictions
        y_test_data: numpy array of y_test data

    Out:
    table: a pandas df of precision, recall, f1, and accuracy scores for your
    ↴model
    """
    accuracy = accuracy_score(y_test_data, preds)
    precision = precision_score(y_test_data, preds)
    recall = recall_score(y_test_data, preds)
    f1 = f1_score(y_test_data, preds)
```

```

table = pd.DataFrame({'model': [model_name],
                      'precision': [precision],
                      'recall': [recall],
                      'F1': [f1],
                      'accuracy': [accuracy]
                     })

return table

```

1. Use the `get_test_scores()` function to generate the scores on the test data. Assign the results to `rf_test_scores`.
2. Call `rf_test_scores` to output the results.

RF test results

```
[210]: # Get scores on test data
rf_test_scores = get_test_scores('RF test', rf_preds, y_test)
results = pd.concat([results, rf_test_scores], axis=0)
results
```

```
[210]:    model  precision      recall        F1  accuracy
0    RF CV    0.674760  0.757467  0.713618  0.680151
0   RF test    0.669735  0.769757  0.716271  0.679004
```

Question: How do your test results compare to your validation results?

The test results are very similar to the cross-validation (CV) results, indicating that the model generalizes well to unseen data.

CV F1: 0.7133

Test F1: 0.7163

The differences across all metrics (precision, recall, accuracy, F1) are very small — all within ~0.005. This suggests that the Random Forest model is neither overfitting nor underfitting, and that the cross-validation process provided a reliable estimate of real-world performance.

Overall, the model's performance on the test set confirms that it maintains consistent predictive quality outside the training folds.

XGBoost Try to improve your scores using an XGBoost model.

1. Instantiate the XGBoost classifier `xgb` and set `objective='binary:logistic'`. Also set the random state.
2. Create a dictionary `cv_params` of the following hyperparameters and their corresponding values to tune:
 - `max_depth`
 - `min_child_weight`
 - `learning_rate`

- `n_estimators`
3. Define a set `scoring` of scoring metrics for grid search to capture (precision, recall, F1 score, and accuracy).
 4. Instantiate the `GridSearchCV` object `xgb1`. Pass to it as arguments:
 - `estimator=xgb`
 - `param_grid=cv_params`
 - `scoring=scoring`
 - `cv`: define the number of cross-validation folds you want (`cv=_`)
 - `refit`: indicate which evaluation metric you want to use to select the model (`refit='f1'`)

```
[211]: # 1. Instantiate the XGBoost classifier
xgb = XGBClassifier(objective='binary:logistic', random_state=0)

# 2. Create a dictionary of hyperparameters to tune
# Note that this example only contains 1 value for each parameter for
# simplicity,
# but you should assign a dictionary with ranges of values
cv_params = {'learning_rate': [0.1],
             'max_depth': [8],
             'min_child_weight': [2],
             'n_estimators': [500]
            }

# 3. Define a list of scoring metrics to capture
scoring = ['accuracy', 'precision', 'recall', 'f1']

# 4. Instantiate the GridSearchCV object
xgb1 = GridSearchCV(xgb, cv_params, scoring=scoring, cv=4, refit='f1')
```

Now fit the model to the `X_train` and `y_train` data.

```
[212]: %%time
#==> ENTER YOUR CODE HERE
xgb1.fit(X_train, y_train)
```

CPU times: user 23.1 s, sys: 152 ms, total: 23.2 s
 Wall time: 12.1 s

```
[212]: GridSearchCV(cv=4,
                    estimator=XGBClassifier(base_score=None, booster=None,
                                            callbacks=None, colsample_bylevel=None,
                                            colsample_bynode=None,
                                            colsample_bytree=None, device=None,
                                            early_stopping_rounds=None,
                                            enable_categorical=False, eval_metric=None,
                                            feature_types=None, gamma=None,
```

```

grow_policy=None, importance_type=None,
interaction_constraints=None,
learning_rate=None, ...
max_delta_step=None, max_depth=None,
max_leaves=None, min_child_weight=None,
missing=nan, monotone_constraints=None,
multi_strategy=None, n_estimators=None,
n_jobs=None, num_parallel_tree=None,
random_state=0, ...),
param_grid={'learning_rate': [0.1], 'max_depth': [8],
            'min_child_weight': [2], 'n_estimators': [500]},
refit='f1', scoring=['accuracy', 'precision', 'recall', 'f1'])

```

Get the best score from this model.

```
[213]: # Examine best score
print("Best CV F1 Score:", xgb1.best_score_)
```

Best CV F1 Score: 0.6974565593334565

And the best parameters.

```
[214]: # Examine best parameters
print("Best Parameters:", xgb1.best_params_)
```

Best Parameters: {'learning_rate': 0.1, 'max_depth': 8, 'min_child_weight': 2, 'n_estimators': 500}

XGB CV Results Use the `make_results()` function to output all of the scores of your model. Note that it accepts three arguments.

```
[215]: # Call 'make_results()' on the GridSearch object
xgb1_cv_results = make_results('XGB CV', xgb1, 'f1')
results = pd.concat([results, xgb1_cv_results], axis=0)
results
```

	model	precision	recall	F1	accuracy
0	RF CV	0.674760	0.757467	0.713618	0.680151
0	RF test	0.669735	0.769757	0.716271	0.679004
0	XGB CV	0.671738	0.725420	0.697457	0.668768

Use your model to predict on the test data. Assign the results to a variable called `xgb_preds`.

HINT

You cannot call `predict()` on the `GridSearchCV` object directly. You must call it on the `best_estimator_`.

```
[216]: # Get scores on test data
xgb_preds = xgb1.best_estimator_.predict(X_test)
```

XGB test results

1. Use the `get_test_scores()` function to generate the scores on the test data. Assign the results to `xgb_test_scores`.
2. Call `xgb_test_scores` to output the results.

```
[217]: # Get scores on test data
xgb_test_scores = get_test_scores('XGB test', xgb_preds, y_test)
results = pd.concat([results, xgb_test_scores], axis=0)
results
```

```
[217]:      model  precision    recall        F1  accuracy
0    RF CV    0.674760  0.757467  0.713618  0.680151
0   RF test    0.669735  0.769757  0.716271  0.679004
0   XGB CV    0.671738  0.725420  0.697457  0.668768
0   XGB test    0.676471  0.758556  0.715166  0.681952
```

Question: Compare these scores to the random forest test scores. What do you notice? Which model would you choose?


The Random Forest and XGBoost models have nearly identical performance. XGBoost slightly outperforms Random Forest in precision and overall accuracy, while Random Forest has a marginally higher recall. Since both models generalize well, I would select XGBoost for deployment because of its efficiency and slightly better balance of metrics on unseen data.

Plot a confusion matrix of the model's predictions on the test data.

```
[218]: # Generate array of values for confusion matrix
cm = confusion_matrix(y_test, rf_preds, labels=rf1.classes_)

# Plot confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                               display_labels=rf1.classes_,
                               )
disp.plot(values_format='1')
```

```
[218]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7aba26cac3a0>
```




Question: What type of errors are more common for your model?

The model makes more False Positive errors than False Negatives. This means it's more likely to predict that a passenger will be generous when they aren't. In other words, the model slightly overpredicts generosity. For a taxi driver, this might mean sometimes expecting a higher tip than they actually receive.

Feature importance Use the `feature_importances_` attribute of the best estimator object to inspect the features of your final model. You can then sort them and plot the most important ones.

```
[219]: importances = rf1.best_estimator_.feature_importances_
rf_importances = pd.Series(importances, index=X_test.columns)
rf_importances = rf_importances.sort_values(ascending=False)[:15]

fig, ax = plt.subplots(figsize=(8,5))
rf_importances.plot.bar(ax=ax)
ax.set_title('Feature importances')
ax.set_ylabel('Mean decrease in impurity')
fig.tight_layout();
```



4.4 PACE: Execute

Consider the questions in your PACE Strategy Document to reflect on the Execute stage.

4.4.1 Task 4. Conclusion

In this step, use the results of the models above to formulate a conclusion. Consider the following questions:

1. Would you recommend using this model? Why or why not?

Yes, I would recommend using this model — particularly the XGBoost model. Both Random Forest and XGBoost achieved similar F1 and accuracy scores (around 0.71 F1 and 0.68 accuracy), showing good consistency between training and testing results. The XGBoost model slightly outperformed Random Forest in precision and overall accuracy, indicating that it generalizes well and avoids overfitting. While performance isn't perfect, it's reliable enough to make useful predictions about passenger generosity.

2. What was your model doing? Can you explain how it was making predictions?

The model predicts whether a passenger is generous (high tip) or not generous based on trip-related features. The most important features influencing predictions were:

VendorID_2

predicted_fare

mean_duration and mean_distance

passenger_count These suggest that longer or more expensive rides, or certain vendors, are correlated with higher tipping behavior. Essentially, the model learns patterns between trip length, cost, and timing to estimate generosity probability.

3. Are there new features that you can engineer that might improve model performance?

Yes — performance could likely be improved with richer or more contextual features, such as:

Weather conditions (e.g., rain or temperature at pickup time — might affect tipping behavior).

Traffic or congestion level (trip delays might influence satisfaction and tips).

Passenger type (e.g., airport pickups vs. local rides).

Time since last trip or driver experience metrics (if available). These features could provide the model with more insight into passenger context and behavior.

4. What features would you want to have that would likely improve the performance of your model?

I would like to include:

Real-time trip context: weather, time of day, and location demographics.

Driver and passenger profiles: previous tip behavior, ratings, or loyalty indicators.

Trip purpose or destination type: e.g., hotel, airport, or residential area. These would allow the model to learn more nuanced patterns behind why certain trips yield higher tips.

Remember, sometimes your data simply will not be predictive of your chosen target. This is common. Machine learning is a powerful tool, but it is not magic. If your data does not contain predictive signal, even the most complex algorithm will not be able to deliver consistent and accurate predictions. Do not be afraid to draw this conclusion. Even if you cannot use the model to make strong predictions, was the work done in vain? Consider any insights that you could report back to stakeholders.

Congratulations! You've completed this lab. However, you may not notice a green check mark next to this item on Coursera's platform. Please continue your progress regardless of the check mark. Just click on the "save" icon at the top of this notebook to ensure your work has been logged.